Foundations

Michael Barz

October 14, 2017

Contents

Note: log **(or** lg**) denotes the binary logarithm** $(\log_2)$**, not** $\log_{10}$ **or** $\log_e$.

1 Algorithms

This is an interesting read, but nothing really noteworthy.

2 Getting Started

2.1 Insertion Sort

Insertion sort will be our first algorithm, and it will solve the sorting problem:

> **Sorting Problem**
>
> **Input:** A sequence of n numbers $(a_1, a_2, \ldots, a_n)$.
> **Output:** A permutation $(a'_1, \ldots, a'_n)$ of the input sequence such that $a'_1 \le a'_2 \le \cdots \le a'_n$.

The numbers sorted are called *keys*.

Insertion sort:

```
1  for j=2 to A.length
2        key = A[j]
3        // Insert A[j] into the sorted sequence A[1..j-1]
4        i = j-1
5        while i > 0 and A[i] > key
6                A[i+1] = A[i]
7                i = i-1
8        A[i+1] = key
```

Insertion sort works by sorting subsections of the list.

First, it sorts the first 2 elements, then the first 3, ...

It does this by taking element j of the list. Then, it compares it $A[j - 1], A[j - 2], \ldots$ until it finds an element smaller than $A[j]$. Once it does that, it just puts $A[j]$ in that elements spot. While it's doing this, if an element is larger than $A[j]$, it moves it up one in the list.

> **Loop Invariant**
>
> A **loop invariant** is something which is true at the start of every itera-
> tion of a loop. Formally, it satisfies two properties:
>
> 1. **Initalization:** It is true at the start of the loop's first iteration.
>
> 2. **Maintenance:** If it is true at the start of an iteration of the loop,
>   it is true at the start of the next iteration.
>
> Informally, it also has a third property: **termination.** This means that,
> when the loop terminates, the invariant gives us a useful property which
> helps prove the algorithm correct.

2.1.1 Exercises

Exercise 2.1-3. Consider the **searching problem:**

> **Searching Problem**
>
> **Input:** A sequence of n numbers $A = (a_1, \ldots, a_n)$ and a value v.
> **Output:** The smallest index i such that $v = A[i]$ or the special value
> -1 if v does not appear in A.

Write pseudocode for *linear search,* which scans through the sequence, look-
ing for v. Using a loop invariant, prove that your algorithm is correct.

   Solution. I claim the algorithm

```
1  for i=1 to A.length
2          if A[i] == v
3                  return i
4  return -1
```

Now, we will prove it works by using the following loop invariant: At the
start of each iteration, the sub-sequence $(a_1, \ldots, a_{i-1})$ will not contain v.

To prove that the suggested loop invariant holds, note that at the start of
the loops iteration, the sub-sequence described will be the empty sub-sequence,
which trivially does not contain v.

To prove the claimed invariant has maintanence, note that, if it is true at
the start of an iteration, then $(a_1, \ldots, a_{i-1})$ will not contain v. If a_i is v, then
the algorithm will terminate. However, if the loop runs again, then $a_i \neq v$, and
thus $(a_1, \ldots, a_{i-1}, a_i)$ does not contain v.

Thus, the loop will either terminate the algorithm by returning the index,
or the loop will terminate. If the loop terminates, we will have that $(a_1, \ldots, a_n)$
will not contain v (by the invariant, if it got to the iteration where $i = n$, then
$(a_1, \ldots, a_{n-1})$ will not contain v, and if the loop went through $i = n$ without
returning anything, then we don't have $a_n = v$), and thus we return -1, as
desired.

If the loop does terminate, then it returns the smallest i such that $a_i = v$. To see this, assume that the loop halts at iteration i and returns i. Then, $(a_1, \ldots, a_{i-1})$ will not contain v by the loop invariant, and thus there will be no smaller index k such that $a_k = v$.

Exercise 2.1-4. Consider the binary addition problem:

> ### Binary Addition Problem
>
> **Input:** Two n-element sequences A and B. Each element of A and B will be either 0 or 1, and A and B will represent two binary integers.
> **Output:** The sum of the two integers that A and B represent, represented similarly as an $n+1$-element sequence C.

Find (with proof) an algorithm which solves this problem.

Solution. Binary addition is fairly easy, and the only big issue is carrying. In pseduocode, I'm going to assume that C starts out as a sequence of $n+1$ zeroes, as I don't know how I initalize objects with this language.

```
1  for i=1 to A.length
2          C[i] = C[i] + A[i] + B[i]
3          if C[i] == 2
4                  C[i+1] = 1
5                  C[i] = 0
6          if C[i] == 3
7                  C[i+1] = 1
8                  C[i] = 1
9  return C
```

Note: Numbers are backwards–that is, 110 would be $[0, 1, 1]$, not $[1, 1, 0]$. It's easier for me to implement and prove it works using that, and the only difference is to humans (which you can change by reversing the sequence before outputting/after inputting).

Now, a proof that it works.

Let A_k denote the number represented by the first k digits of A. Similarly define B_k and C_k.

First, we establish a loop invariant: Before the i^{th} iteration of this loop, C represents the sum of the numbers represented by the first $i-1$ digits of A and B.

For initalization, note that before the first iteration, the first 0 digits of A and B aren't anything, and thus their sum is 0–the empty sum–and C represents 0.

Now, proving maintenance. If the property holds on iteration i, then we will prove it holds on iteration $i+1$.

At the start of iteration i, we have that $C_i = A_{i-1} + B_{i-1}$. Also note that

$$A_i = A_{i-1} + 2^{i-1} A[i],$$

4

$$B_i = B_{i-1} + 2^{i-1}B[i].$$

Thus,

$$
\begin{aligned}
A_i + B_i &= A_{i-1} + B_{i-1} + 2^{i-1}(A[i] + B[i]) \\
&= C_i + 2^{i-1}(A[i] + B[i]) \\
&= C_{i-1} + 2^{i-1}(A[i] + B[i] + C[i]).
\end{aligned}
$$

Now, line 2 will set $C[i]$ to $A[i] + B[i] + C[i]$.

If $C[i]$ is 0 or 1, no worries. Then, we just leave $C[i+1]$ as 0, and then, as

$$A_i + B_i = C_{i-1} + 2^{i-1}C[i],$$

the value of $C[i]$ is exactly what the i^{th} digit of the sum $A_i + B_i$ should be. Thus, C_{i+1} equals $A_i + B_i$, as desired.

If $C[i] = 2$, then there's a carry–we have

$$A_i + B_i = C_{i-1} + 2^i,$$

and thus need to place 10 in front of C_{i-1}'s binary representation (equivalent to adding 0 and then 1 to our representation) to get A_i+B_i. This is accomplished by the if statement on line 3.

Similarly, if $C[i] = 3$, we need to place 11 in front of C_{i-1}, which we do with the if statement on line 6.

As $C[i]$ cannot exceed 3, we've covered every case. Thus, maintenance holds.

Anyways, using our loop invariant, at the end of the loop we have that C contains the sum of $A_n + B_n$, which is what we want.

2.2 Algorithm Analysis

Let's look at the insertion sort from before. Except this time, let's analyze how long it takes to run.

By 'cost,' we mean the length of time it takes to execute that line; by 'times,' we mean how many times that line will execute.

There is a while loop in this code–we will denote by t_j the number of times that while loop executes for the given j value. For example, if it executes once for $j = 2$, then we'd say $t_2 = 1$. (This notation is a little different than the books–specifically, they define t_j to be one more than this, but I prefer it.)

We will also denote by c_i the amount of time it takes for line i to execute.

Now, let's look back at that insertion sort:

```
1  for j=2 to A.length
2          key = A[j]
3          // Insert A[j] into the sorted sequence A[1..j-1]
4          i = j-1
5          while i > 0 and A[i] > key
6                  A[i+1] = A[i]
7                  i = i-1
8          A[i+1] = key
```

Line 1 will run n times–it will run for $2, 3, \ldots, n$, and then run for $n + 1$, at which point it will stop.

Lines 2 through 4 will all run $n - 1$ times each–once for each iteration of the loop. Line 8 will also run $n - 1$ times.

Lines 5 through 7 are a little trickier. Line 5 will execute $\sum_{j=2}^{n}(t_j + 1)$ times, and lines 6 and 7 will each run $\sum_{j-2}^{n} t_j$ times.

Now, the total time cost, $T(n)$, of our algorithm is:

$$T(n) = nc_1 + (n-1)(c_2 + c_4 + c_8) + (\sum_{j=2}^{n}(t_j + 1))c_6 + (\sum_{j=2}^{n} t_j)(c_5 + c_7).$$

Now, the best possible scenairo for insertion sort happens when the list given to us is already sorted. Here, we have $t_j = 0$ for all relevant j, and thus the running time becomes

$$T(n) = nc_1 + (n-1)(c_2 + c_4 + c_8) + (\sum_{j=2}^{n}(1))c_6 + (\sum_{j=2}^{n} 0)(c_5 + c_7)$$
$$= n(c_1 + c_2 + c_4 + c_8 + c_6) - (c_2 + c_4 + c_8 + c_6)$$
$$= an + b,$$

for some constants a and b. Thus, in the best case, $T(n)$ grows linearly.

In the worst case, when the array is in reverse order, we have $t_j = j - 1$, as we need to put the j^{th} element at the beginning of the list, and we look one at a time. Thus, we have

$$T(n) = n(c_1 + c_2 + c_4 + c_8) - (c_2 + c_4 + c_8) + (n(n+1)/2 - 1)c_6 + (n(n-1)/2)(c_5 + c_7)$$
$$= an^2 + bn + c,$$

for some constants a, b, c. Thus, in the worst case, $T(n)$ grows quadratically.

2.2.1 Exercises

Exercise 2.2-2. Consider sorting n numbers stored in array A by first finding the smallest element of A and exchanging it with the element in $A[1]$. Then find the second smallest element and exchange it with $A[2]$, and so on. Write pseudocode for this algorithm, known as a *selection sort*. Give the best and worst case running times (using Θ notation).

Solution. First, we describe it's implementation (given some sequence A):

```
1  for i=1 to A.length-1
2          smallest = i
3          for j=i+1 to A.length
4                  if A[j] < A[smallest]
5                          smallest = j
```

```
6              if smallest != i // don't␣switch␣unless␣we␣need␣to
7 ␣␣␣␣␣␣␣␣␣␣␣␣␣␣␣␣␣temp␣=␣A[smallest]
8 ␣␣␣␣␣␣␣␣␣␣␣␣␣␣␣␣␣A[smallest]␣=␣A[i]
9 ␣␣␣␣␣␣␣␣␣␣␣␣␣␣␣␣␣A[i]␣=␣temp
```

Note that, with this implementation, the best and worst case running times are both $\Theta(n^2)$. This is because the only place where we can save time is skipping stuff in that if statement, which can only save linear time, whereas the inner for loop has a quadratic running time.

Exercise 2.2-3. What are the average, best, and worse case running times of the linear search from earlier?

Solution. The best case, when $A[1] = v$, gives us a constant running time. The worst case, when $A[n] = v$, gives us linear running time.

If the item occurs once in the sequence, then we expect it to be at position

$$\frac{1 + 2 + \cdots + n}{n} = \frac{n(n+1)/2}{n} = \frac{n+1}{2},$$

and thus the running time is linear.

2.3   Merge Sort

Merge sort is a divide and conquer solution to the sorting problem. It has three steps:

1. Divide the n-element sequence into two $n/2$-element sequence.

2. Sort the sub-sequences using merge sort.

3. Combine the sorted sub-sequences to make a sorted sequences.

We'll define `merge(A, p, q, r)` to take a sequence of numbers A such that $A[p..q]$ and $A[q+1..r]$ are sorted, and then merges them to form a single sorted sub-sequence, which will be stored in $A[p..r]$. Our algorithm for `merge` takes $\Theta(n)$ time, where $n = r - p + 1$ is the size of the sub-sequence.

```
1  n_1 = q-p+1
2  n_2 = r-q
3  let L[1..n_1 + 1] and R[1..n_2+1] be new arrays
4
5  for i=1 to n_1
6          L[i] = A[i+p-1]
7  for i=1 to n_2
8          R[i] = A[i+q]
9  L[n_1+1] = \infty
10 R[n_2+1] = \infty
11
12 i = 1
```

```
13  j = 1
14  for k = p to r
15          if L[i] \leq R[j]
16                  A[k] = L[i]
17                  i = i+1
18          else
19                  A[k] = R[j]
20                  j = j+1
```

The proof that this works is kinda boring, so I'll skip it.

Now, merge sort!

`merge-sort(A, p, r)` accepts a sequence A, an two indices p and r, and sorts $A[p..r]$.

```
1  if p < r
2          q = [(p+r)/2]
3          merge-sort(A, p, q)
4          merge-sort(A, q+1, r)
5          merge(A, p, q, r)
```

To sort the entire sequence A, we call `merge-sort(A, 1, A.length)`.

Also, $[x]$ denotes the floor of x.

2.4   Running Time of Merge Sort

Let's compute $T(n)$, the worst case running time of merge sort on a set of n numbers. Merge sort on $n = 1$ takes constant time, otherwise it takes $2T(n/2)$ time to sort the two subsequences, and then some linear amount of time to merge. Thus,

$$T(n) = \begin{cases} c & n = 1 \\ 2T([n/2]) + \Theta(n) & n > 1 \end{cases}$$

Rigorously solving this requires actually defining Θ notation, but intuitively it's $\Theta(n \lg n)$.

2.4.1   Exercises

Exercise 2.3-2. Rewrite the `merge` procedure so that it doesn't use infinity.

Solution. See `merge_sort.py`, the python implementation of the algorithm.

Exercise 2.3-3. Use mathematical induction to show that when n is an exact power of 2, the solution to the recurrence

$$T(n) = \begin{cases} 2 & \text{if } n = 2, \\ 2T(n/2) + n & \text{if } n = 2^k, k > 1 \end{cases}$$

is $T(n) = n \lg n$.

Solution. If $n = 2$, then $T(n) = 2 = 2\lg 2$ trivially. Now, the inductive step.

Assume that $T(2^k) = k2^k$. Then,

$$T(2^{k+1}) = 2T(2^k) + 2^{k+1} = 2k2^k + 2(2^k) = 2^k(2k+2) = 2^{k+1}(k+1),$$

as desired.

Exercise 2.3-4. Recursively, we can define insertion sort to sort $A[1..n]$ by sorting $A[1..n-1]$ and then insert $A[n]$ into the sorted array $A[1..n-1]$. Write a recurrence for the running time of this algorithm.

Solution. Let $T(n)$ be the running time. Then,

$$T(n) = \begin{cases} c & \text{if } n = 1 \\ T(n-1) + \Theta(n) & \text{otherwise} \end{cases}$$

where c is some constant.

Exercise 2.3-5. Consider the searching problem. If A is sorted, we can use tthe binary search algorithm, which checks the midpoint and then halves. Implement it, prove it works, then argue that the worst case running time is $\Theta(\lg n)$.

Solution. I implement the `binary_search(A, v, n, m)` function. A is a sequence, v is the value, and n, m are the start and end of the search interval. To search, call `binary_search(A, v, 1, A.length)`.

```
1  if n > m
2          return -1
3  else if A[n] == v
4          return n
5  else if n == m
6          return -1
7
8  x =   ceil((n+m) / 2)
9
10 if A[x] == v
11         return x
12 else if A[x] > v
13         return binary_search(A, v, n, x-1)
14 else
15         return binary_search(A, v, x+1, m)
```

Now, we will prove it works. To prove it works, note that if it is called, then either v is nowhere in the list, or it is in $A[n..m]$. This can be proven by induction

easily. For the base case, when we call it on $n = 1$ and $m = A.length - 1$, its trivially true.

Now, assume it was called on n and m. Then, we check the midpoint of the interval. If we don't find it by checking the midpoint x, then it's either on $A[n..x-1]$ or $A[x+1..m]$ or not in A. If it's on $A[x] > v$, then it must be on $A[n..x-1]$ if it's in A by our list. Otherwise, it must be on $A[x+1..m]$ if it's on our list. Thus, at the start of each method call (assuming we started with $n = 1$ and $m = A.length - 1$) the invariant remains true.

Now, the size of the interval $A[n..m]$ will continually decrease. Eventually, if we have not found v earlier, we will get to an interval containing only one point. In this case, we explicitly check if $A[n] = v$, and proceed accordingly.

Anyways, now that we know it works, let's find the time complexity.

The method takes some constant time, say c, to run on a single iteration. Say the first 6 lines take c' time to run, where c' is constant. Thus, we have

$$T(n) = \begin{cases} c' & \text{if n=1} \\ T(n/2) + c \end{cases}$$

which is $\Theta(\lg n)$, as it will take $\lg n$ iterations to go stop this recursion, and other then that we only add constants.

Exercise 2.3-7. Describe a $\Theta(n \lg n)$ algorithm that, given a set S of n integers and another integer x, determines whether or not there exist two elements in S whose sum is exactly x.

Solution. We can do this. Consider the algorithm `sum(A, x)`

1. Merge sort the list.

2. Iterate over the first $n-1$ elements of S, and binary search to see if $x-i \in S$ for each $i \in S$. Use binary search for this.

Merge sort is $\Theta(n \lg n)$ and binary search is $\Theta(\lg n)$. As the binary search will be repeated n times, the overall time complexity is $\Theta(n \lg n)$ for the second step, and thus the total algorithm has a time complexity of $\Theta(n \lg n)$.

2.5 Problems

Problem 2-1. Although merge sort runs in $\Theta(n \lg n)$ worst-case time and insertion sort in $\Theta(n^2)$, the constant factors in insertion sort makes it faster when you have sufficiently small lists. Consider a modification to merge sort in which n/k sublists of length k are sorted using insertion sort and then merged using the standard merging mechanism, where k is a value to be determined.

1. Show that insertion sort can search all n/k sublists in $\Theta(nk)$ worst-case time.

2. Show how to merge the sublists in $\Theta(n \lg(n/k))$ worst-case time.

3. Given the modified algorithm runs in $\Theta(nk + n \lg(n/k))$ worst-case time, what is the largest value of k as a function of n for which the modified algorithm has the same running time as standard merge sort, in terms of Θ notation?

4. How do we choose k in practice?

Solution. For part (a), consider that insertion sort has a worst-case time of $\Theta(k^2)$ for each sublist. Thus, n/k sublists with $\Theta(k^2)$ worst time is a total time of $\Theta(n/k \cdot k^2) = \Theta(nk)$.

Problem 2-2. Bubblesort is a sorting algorithm. It works by repeatedly swapping adajcent elements that are out of order. Prove bubblesort is correct. What is the worst case running time?

Solution. The code for bubble sort is included below.

```
1  for i = 1 to A.length - 1
2          for j = A.length downto i+1
3                  if A[j] < A[j-1]
4                          exchange A[j] with A[j-1]
```

To prove it works, we will first establish some loop invariants.

Consider that inner for loop. At the start of each iteration, the smallest element of $A[j..A.length]$ will be $A[j]$.

At the start of the first iteration, $j = A.length$ so the claim is trivially true.

Now, assume the claim is true at the start of iteration j. Then, if $A[j]$ is smaller than $A[j-1]$, we will swap those elements so now $A[j-1]$ will be the msaller of the two. Note that, by hypothesis, $A[j]$ is the smallest element of $A[j..A.length]$, and thus after that swap (if it happens), $A[j-1]$ will be the smallest element of $A[j-1..A.length]$, as desired.

Thus, after that inner for loop executes, $A[i]$ will be the smallest element of $A[i..A.length]$.

Now, a loop invariant for the main for loop. Before iteration i of the loop, the smallest $i-1$ elements of A will be those in $A[1..i-1]$, and $A[1..i-1]$ will be in non-decreasing order. For the first iteration, this claim is trivially true as $A[1..0]$ is an empty list.

Now, assume that it's true before iteration i. By hypothesis, we have that $A[i-1] \leq x$ for any $x \in A[i..A.length]$. In particular, $A[i-1] \leq m$, where m is the minimum element of $A[i..A.length]$.

By the inner for loop's invariant, after it terminates, we will have that $A[i]$ is the minimum element of $A[i..A.length]$. Thus, we will have $A[i-1] \leq A[i]$ (as $A[i] = m$ as disucssed above), and so $A[1..i]$ will be in non-creasing order. Similarly, the smallest i elements of A will be those in $A[1..i]$, as desired. Thus, the loop invariant holds.

By this loop invariant, after iteration $A.length - 1$ of the loop, we will have $A[1..A.length - 1]$ is sorted and that $A[A.length - 1] \leq A[A.length]$, implying that A is sorted. This proves correctness.

Now, to find the worst case running time.

Note that line 1 takes constant time to run, and it will run n times.

Line 2 also takes constant time, and it will run $(n + (n - 1) + \cdots + 2) = (n(n + 1)/2 - 1) = (n^2 + n - 2)/2$ times.

Lines 3 and 4 both take constant time. Line 3 will run $(n^2 + n)/2$ times,, and in the worst case line 4 will also run $(n^2 + n)/2$ times.

Thus, the total running time will be $\Theta(n^2)$, the same as insertion sort.

Problem 2-3. Horner's rule is a rule for evaluating a polynomial, which states that

$$P(x) = \sum_{k=0}^{n} a_k x^k$$
$$= a_0 + x(a_1 + x(a_2 + \cdots + x(a_{n-1} x a_n) \ldots)).$$

Using Θ notation, find the running time of using Horner's rule to evaluate a polynomial. Compare this to the naive algorithm which computes $a_0 + a_1 x + \cdots + a_n x^n$ by finding each term from scratch.

Solution. We can implement Horner's rule like so:

```
1  y = 0
2  for i = n downto 0
3          y = a_i + xy
```

To find it's running time, note that line 1 takes constant time and runs once. Line 2 will run $n + 2$ times (and takes constant time) and line 3 will run $n + 1$ times (also taking constant time). Thus, the algorithm is $\Theta(n)$.

For the naive algorithm, that depends on how exponentiation is implemented. If it's implemented as an algorithm in terms of multiplication with $\Theta(n)$ complexity, then the naive algorithm is $\Theta(n^2)$. If exponentiation can be done in constant time for arbitrary numbers, then its still $\Theta(n)$ (although this is very unlikely).

Problem 2-4. Let $A[1..n]$ be an array of n distinct numbers. If $i < j$ and $A[i] > A[j]$, then the pair (i, j) is called an inversion of A.

What array whose elements are from the set $\{1, 2, \ldots, n\}$ has the most inversions, and how many does it have?

What is the relationship between the running time of insertion sort and the number of inversions in the input sequence?

Give an algorithm that determines the number of inversions in any permutation on n elements in $\Theta(n \lg n)$ worst-case time.

Solution.

The array $(n, n-1, n-2, \ldots, 1)$ has the most inversions of any permutation, with a whopping $\binom{n}{2}$ inversions (given any $0 < i < j < n+1$, we have that (i, j) is an inversion). To prove this, we use induction on n. For $n = 1$, this is trivial.

Now, suppose that $(k, k-1, \ldots, 1)$ is the permutation with the most inversions of any permutation of $\{1, \ldots, k\}$. Then, consider a permutation of $\{1, \ldots, k+1\}$.

Suppose some permutation of $\{1, \ldots, k+1\}$ had more inversions than $(k+1, \ldots, 1)$. Suppose P was our permutation, and that $k+1$ was placed at index i. If $i = 1$, the the number of inversions is k plus the number of inversions of $A[2..k+1]$, which (by induction) is maximixed when $A[2..k+1] = (k, k-1, \ldots, 1)$.

Now, suppose that $i > 1$. Then, $A[1] < A[i]$, so that can't be an inversion. But every possible (i, j) pair (with $0 < i < j < k+1$) was an inversion in $(k+1, \ldots, 1)$ so this permutation can't possibly have more inversions, since we've found one pair that isn't an inversion.

3 Growth of Functions

We can compute the exact running time of an algorithm as a function of the input size, but that's usually not worth the effort. Instead, we study the *asymptopic* efficiency of algorithms. That is, we are concerned with how the running time of an algorithm increases as the size of the input increases without beyond.

3.1 Asymptotic Notation

Wheenver we define a function representing the running time of an algorithm, we assume it is a function whose domain is $\mathbb{N}$ (like any sensible person would, this book defines 0 to be a natural number). Sometimes we abuse notation and extend it to real numbers or restrict it to a subset of the naturals.

3.1.1 Θ-notation

Θ-notation

Given a function $g(n)$, we define

$$\Theta(g(n)) = \{f(n) : \text{there exists } c_1, c_2, n_0 > 0$$
$$\text{such that } 0 \leq c_1 g(n) \leq f(n) \leq c_2 g(n) \text{for all } n \geq n_0)\}.$$

If $f(n) \in \Theta(g(n))$, we say that g is an *asymptotically tight bound* for f.

The above definition just means that $\Theta(g(n))$ is just the set of all functions that grow at the same rate (up to some constant factor) as $g(n)$ for sufficiently large n. The condition that $f(n) \leq c_2 g(n)$ is fairly obvious, as if it could get bigger than $c_2 g(n)$ for all c_2 it would definitely grow faster; the $c_1 g(n) \leq f(n)$ is so it isn't slower than $g(n)$.

The definition requires that f be *asymptotically nonnegative,* which means that f is nonnegative for sufficiently large n. In order for $\Theta(g(n)) \neq \emptyset$, we must have that g be asymptotically nonnegative as well.

Instead of writing $f(n) \in \Theta(g(n))$, we usually write $f(n) = \Theta(g(n))$.

3.1.2 O-notation

Θ-notation asymptotically bounds a function from above and below. O-notation only bounds it from above.

O-notation

Let $g(n)$ be a function. Then we define

$$O(g(n)) = \{f(n) : \text{there exists } c, n_0 > 0$$
$$\text{such that } 0 \leq f(n) \leq cg(n) \text{ for all } n \geq n_0\}.$$

3.1.3 Ω-notation

Ω-notation gibes us asymptotic lower bounds.

ω-notation

Let $g(n)$ be a function. Then we define

$$\Omega(g(n)) = \{f(n) : \text{there exists } c, n_0 > 0$$
$$\text{such that } 0 \leq cg(n) \leq g(n) \text{ for all } n \geq n_0\}.$$

When we say that the running time of an algorithm is $\Omega(g(n))$, we mean that *no matter what particular input of size n is chosen,* the running time on that input is at least a constant times $g(n)$ (for sufficiently large n).

3.1.4 Notes on Notation

Now that these notations are all defined, we can state and prove the following simple theorem:

Theorem 3.1

For any two functions $f(n)$ and $g(n)$, we have $f(n) = \Theta(g(n))$ if and only if $f(n) = O(g(n))$ and $f(n) = \omega(g(n))$.

This proof is trivial; just look back at the definitions.

This theorem is mainly useful for proving that $f(n) = \Theta(g(n))$ by first proving an upper and lower bound for f.

There's one more important thing to mention: Sometimes, we use $\Theta(g(n))$ to represent a function $f(n) \in \Theta(g(n))$ that we don't want to write out. For example, we may write

$$2n^2 + 3n + 1 = 2n^2 + \Theta(n).$$

This can help eliminate clutter in our algebra; for example, in chapter 2 we found that merge sort had worst-case running time

$$T(n) = 2T(n/2) + \Theta(n).$$

As we are only interested in the asymptotic behaviour of $T(n)$, there is no point in specifying exactly which function in $\Theta(n)$ we're referring to.

Sometimes, we have equations like

$$2n^2 + \Theta(n) = \Theta(n^2),$$

where Θ-notation appears on both sides. Here, we just mean that no matter which function in $\Theta(n)$ we choose to substitute in on the left hand side, the left hand side will be a member of $\Theta(n^2)$.

This let's us chain things; for example,

$$an^2 + bn + c = an^2 + \Theta(n) = \Theta(n^2).$$

3.1.5 o-notation

The asymptotic upper bound provided by O-notation may or may not be asymptotically tight; we use o-notation to say that it's not tight.

Formally, we have

---

o-notation

Let $g(n)$ be a function. Then we define

$$o(g(n)) = \{f(n) : \text{for all } c > 0 \text{ there exists a constant } n_0 > 0$$
$$\text{such that } 0 \leq f(n) < cg(n) \text{ for all } n \geq n_0\}.$$

For example, $2n = o(n^2)$ but $n^2 \neq o(n^2)$.

$f(n) = o(g(n))$ is equivalent to stating that

$$\lim_{n \to \infty} \frac{f(n)}{g(n)} = 0.$$

3.1.6 ω-notation

Similarly, ω-notation is to Ω-notation as o-notation is to O-notation.

Formally,

---

o-notation

Let $g(n)$ be a function. Then we define

$$\omega(g(n)) = \{f(n) : \text{for all } c > 0 \text{ there exists a constant } n_0 > 0$$
$$\text{such that } 0 \leq cg(n) < f(n) \text{ for all } n \geq n_0\}.$$

For example, $n^2 = \omega(n)$ but $n^2 \neq \omega(n^2)$.

3.1.7   Exercises

Exercise 3.1-1. Let $f(n)$ and $g(n)$ be asymptotically nonnegative functions. Prove that
$$\max\{f(n), g(n)\} = \Theta(f(n) + g(n)).$$

Solution. As f and g are asymptotically nonnegative, let n_1 be some positive constant such that

$$n \geq n_1 \implies f(n), g(n) \geq 0.$$

Let $h(n) = \max\{f(n), g(n)\} = f(n) + g(n) - |f(n) - g(n)|$.

Now, we wish to prove that there exists some positive constants c and n_0 such that $0 \leq h(n) \leq cf(n) + cg(n)$ for all $n \geq n_0$.

Note that

$$h(n) = f(n) + g(n) - |f(n) - g(n)| \leq f(n) + g(n) - 0 = f(n) + g(n).$$

Now, all that's left to do is find some n_0 such that $n \geq n_0$ implies $0 \leq h(n)$.

Note that, by the triangle inequality,

$$|f(n) - g(n)| \leq |f(n)| + |g(n)|.$$

Recall that, when $n \geq n_1$, we have that $|f(n)| = f(n)$ and $|g(n)| = g(n)|$. Thus,

$$n \geq n_1 \implies |f(n) - g(n)| \leq |f(n)| + |g(n)|$$
$$= f(n) + g(n).$$

Thus,

$$n \geq n_1 \implies f(n) + g(n) - |f(n) - g(n)| \geq 0,$$

and thus

$$0 \le h(n) \le f(n) + g(n)$$

for all $n \ge n_1$. Thus, $c = 1$ and $n_0 = n_1$ work, and $h(n) = \Theta(f(n) + g(n))$.

Exercise 3.1-2. Show that, for any real constants a and b with $b > 0$, we have that
$$(n + a)^b = \Theta(n^b).$$

Solution. First, we show that

$$n + a \le |a|n$$

when $n > |a|$. To see this, simply note that

$$n + a \le n + |a| < 2n.$$

As $x \mapsto x^b$ is montonically increasing for positive inputs, we have that

$$(n + a)^b < (2n)^b = 2^b n^b.$$

Thus,
$$(n + a)^b = O(n^b)$$

by definition (for the positive part of the definition, simply note that when $n + a$ is positive, $(n + a)^b$ is positive too).

Now, we will prove that $(n + a)^b = \Omega(n^b)$.

To see this, note that whenever $n > |a|$, we have

$$0 \le n \le n + a,$$

and thus

$$0 \le n^b \le (n + a)^b,$$

as desired.

Thus, by theorem 3.1, we have $(n + a)^b = \Theta(n^b)$.

Exercise 3.1-4. Is $2^{n+1} = O(2^n)$? How about 2^{2n}?

Solution. $2^{n+1} = 2 \cdot 2^n$ is trivially $O(2^n)$, so let's consider 2^{2n}.

I claim the answer is no. To see this, assume there were some constants n_0 and $c > 0$ such that, for all $n > n_0$ we have

$$2^{2n} \le c2^n.$$

17

Note that, as $2^n > 0$ for all n, we have

$$2^n \leq c$$

for all $n > n_0$. However, this statement is trivially false as c is a constant and $2^n \to \infty$ as $n \to \infty$. Thus, 2^{2n} is not $O(2^n)$.

3.2 Stardard Notations and Common Functions

3.2.1 Logarithms

We use $\lg^k n = (lgn)^k$ (not composition as it is sometimes defined) and $\lg \lg n = \lg(\lg n)$.

An important notational convention the book uses is that $\lg n + k$ should be interpreted as $\lg(n) + k$, not $\lg(n + k)$.

We will say that a function $f(n)$ is *polylogarithmically bounded* if $f(n) = O(\lg^k n)$ for some constant k.

Finally, something math-y. For $x > -1$, we have

$$\frac{x}{1 + x} \leq \ln(1 + x) \leq x,$$

with equality holding if and only if $x = 0$.

3.2.2 Factorials

Factorials are defined as expected, and we have the (very) weak inequality $n! \leq n^n$. *Stirling's approximation* is

$$n! = \sqrt{2\pi n} \left(\frac{n}{e}\right)^n \left(1 + \Theta\left(\frac{1}{n}\right)\right),$$

where e is the base of the natural logarithm. This gives us a tighter bound on the factorial.

3.2.3 Functional Iteration

We will use $f^{(i)}(n)$ to denote f applied iteratively i times.

The function $\lg^* n$, the iterated logarithm, is defined as

$$\lg *n = \min\{i \geq 0 : \lg^{(i)} n \leq 1\}.$$

The iterated logarithm grows *very* slowly. We very rarely encounter n such that $\lg *n > 5$.

3.2.4   Exercises

Exercise 3.2-1. Show that if f and g are montonically increasing functions, then so are $f + g$ and $f \circ g$. Also show that, if f and g are montonically increasing and nonnegative, then $f \cdot g$ is montonically increasing.

Solution.

First, note that we have, for all $n < m$,

$$f(n) < f(m), g(n) < g(m).$$

Thus,

$$f(n) + g(n) < f(m) + g(m),$$

as desired.

For composition, note that

$$g(n) < g(m),$$

and thus

$$f(g(n)) < f(g(m))$$

when $n < m$.

Now, if we assume that $f(n), g(n) \geq 0$ and let $n < m$, then

$$f(n) < f(m),$$

$$f(n)g(n) < f(m)g(n) < f(m)g(m).$$

Note that we implicitly use $f(m) \geq 0$ when we say $f(m)g(n) < f(m)g(m)$.

Exercise 3.2-2. Prove $a^{\log_b c} = c^{\log_b a}$.

Solution.

Simply note that

$$a^{\log_b c} = (b^{\log_b a})^{\log_b c} = (b^{\log_b c})^{\log_b a} = c^{\log_b a}.$$

Exercise 3.2-3. Prove that $n! = o(n^n)$, $n! = \omega(2^n)$, and that $\lg(n!) = \Theta(n \lg n)$.

Solution.

First, we prove that $n! = o(n^n)$. Given some positive constant $c > 0$, note that if $n > \frac{1}{|1-c|}$, then we have

$$\frac{1}{|1-c|} > n \implies |1-c| < \frac{1}{n}$$

$$\implies 1 - c < \frac{1}{n}$$

$$\implies c > \frac{n-1}{n}$$

$$\implies \frac{n-1}{c} < n$$

$$\implies \frac{n-1}{c} \cdot (n-2)! < n(n-2)! < n^{n-1}$$

$$\implies \frac{1}{c} \cdot (n-1)! < n^{n-1}$$

$$\implies n! < cn^n,$$

and thus given any constant $c > 0$ we can find some n_0 such that $n > n_0 \implies n! < cn^n$, as desired.

Now, to prove that $n! = \omega(2^n)$. Just let $n > \min\{2, 2c\}$ to get that

$$c \cdot 2^n = (2c)2^{n-1} < n2^{n-1} < 2^n < n!$$

as $n > 2$.

Now, for the hard one! We wish to show that $\lg(n!) = \Theta(n \lg n)$.

Note that, by Stirling's approximation, there exists some constants c_1 and c_2 such that

$$\sqrt{2\pi n} \left(\frac{n}{e}\right)^n \left(1 + \frac{c_1}{n}\right) \le n! \le \sqrt{2\pi n} \left(\frac{n}{e}\right)^n \left(1 + \frac{c_2}{n}\right)$$

for all $n > n_0$ for some constant n_0.

TODO: finish this

Exercise 3.2-4. Is $\lceil \lg n \rceil!$ polynomially bounded? How about $\lceil \lg \lg n \rceil!$?

Solution. $\lceil \lg n \rceil!$ is not polynomially bounded. To see this, note that, by Stirling's approximation, we have that

$$\lceil \lg n \rceil! = \sqrt{2\pi \lceil \lg n \rceil} \left(\frac{\lceil \lg n \rceil}{e}\right)^{\lceil \lg n \rceil} \left(1 + \Theta\left(\frac{1}{\lceil \lg n \rceil}\right)\right) = \Theta((\lceil \lg n \rceil)^{\lceil \lg n \rceil}).$$

Now, assume that $\lceil \lg n \rceil^{\lceil \lg n \rceil}$ was polynomially bounded. Then, there are some positive constants m, c such that

$\lceil \lg n \rceil^{\lceil \lg n \rceil} < cn^m$

for all $n > n_0$.

Note that
$$(\lg n)^{\lg n} \le \lceil \lg n \rceil^{\lceil \lg n \rceil} < cn^m.$$

Assume that
$$(\lg n)^{\lg n} < cn^m.$$

Then,
$$(\lg n)(\lg \lg n) < (m \lg n) + (m \lg c).$$

Let n be such that $\lg \lg n > m + (m \lg c)/(\lg n)$. Then,

$$(\lg n)(m + (m \lg c)/(\lg n)) = m \lg n + m \lg c < (m \lg n) + (m \lg c),$$

a clear contradiction.

$\lceil \lg \lg n \rceil!$ is, however, polynomially bounded. By Stirling's approximation, we have

$$\lceil \lg \lg n \rceil! = \sqrt{2\pi \lceil \lg \lg n \rceil} \left(\frac{\lceil \lg \lg n \rceil}{e} \right)^{\lceil \lg \lg n \rceil} \left(1 + \Theta \left(\frac{1}{\lceil \lg \lg n \rceil} \right) \right) = \Theta((\lceil \lg \lg n \rceil)^{\lceil \lg \lg n \rceil}).$$

I claim that
$$(\lceil \lg \lg n \rceil)^{\lceil \lg \lg n \rceil} < n$$

for all $n > 1$. To see this, let $m > 1$ and suppose that

$$(\lceil \lg \lg m \rceil)^{\lceil \lg \lg m \rceil} \ge m.$$

Then,
$$(1 + \lg \lg m)^{1 + \lg \lg m} \ge m.$$

Taking the logarithm of both sides, we have that

$$(1 + \lg \lg m)(\lg(1 + \lg \lg m)) \ge \lg m.$$

Note that the LHS grows slower than the RHS, and this isn't true for $m = 1$, and thus it cannot be true for $m > 1$.